# MATRIX-cLBP

Methods-Anchored Trial of Ratio-Indexed eXtracts
in Chronic Low Back Pain

Version 1.2 — NIH Submission Build

2026-03-24

---

3-arm, double-blind, placebo-controlled Phase 2 RCT
N = 525 (175/arm) · 12-week treatment · cLBP indication

Simulation-validated · Chemistry-controlled · Failure-rubric anchored

This document is the NIH submission build (Variant B) derived from the locked v1.1 archive.
All simulation results are deterministic (seed 20260323, 2,000 replications).
SHA-256 hashes confirm reproducibility of all 37 output files.

Robert Long
Screwball7605@aol.com

# Contents


**1 Executive Summary** **4**
- 1.1 Problem . . . 4
- 1.2 Solution . . . 4
- 1.3 Computational Validation . . . 4

**2 Why Existing Studies Fail** **5**
- 2.1 The Five Failure Modes . . . 5
- 2.2 Three-Layer Failure Architecture . . . 5
- 2.3 Audit Scoring Distribution . . . 6
- 2.4 Key Findings from the 25-Paper Audit . . . 6

**3 Replacement Study Design** **7**
- 3.1 Specific Aims . . . 7
  - 3.1.1 Aim 1: Efficacy of Standardised WPE in cLBP . . . 7
  - 3.1.2 Aim 2: Matrix Isolation . . . 7
  - 3.1.3 Aim 3: Chemistry-Control Platform . . . 7
- 3.2 Protocol Summary (SPIRIT 2025 v3) . . . 7
  - 3.2.1 Hypothesis Hierarchy . . . 8
  - 3.2.2 Key Inclusion / Exclusion . . . 8
- 3.3 SAP Summary (v1.0, Locked) . . . 8
- 3.4 Confounder Management . . . 8
- 3.5 Retention Protocol . . . 9

**4 Chemistry and QC Controls** **10**
- 4.1 Product Taxonomy . . . 10
- 4.2 Chemistry Schema . . . 10
- 4.3 Locked Chemistry Decisions . . . 10
- 4.4 Batch Lifecycle . . . 10
- 4.5 QC Simulation Results . . . 11

**5 Simulation and Validation Results** **12**
- 5.1 4-Engine Pipeline . . . 12
- 5.2 Primary Power and Effect-Size Sensitivity . . . 12
- 5.3 Dropout Sensitivity . . . 12
- 5.4 Secondary Endpoint Power . . . 12
- 5.5 Sensitivity Analyses . . . 14
  - 5.5.1 SA-1: Multiple Imputation under MAR . . . 14
  - 5.5.2 SA-2: Tipping-Point Analysis . . . 14
  - 5.5.3 SA-8: Rescue Medication as Time-Varying Covariate . . . 14
- 5.6 Robustness Scenarios . . . 14

**6 Known Limitations and Red-Team Targets** **18**
- 6.1 Attack 1: Functional Unblinding . . . 18
- 6.2 Attack 2: H2 Overinterpretation . . . 18
- 6.3 Attack 3: Opioid Endpoint Underpower . . . 18

6.4 Attack 4: MNAR Tipping-Point Plausibility . . . 18
6.5 Attack 5: Chemistry Transportability . . . 19
6.6 Attack 6: H1 Power Adequacy . . . 19

**7 Funding-Facing Summary** **20**
7.1 Significance . . . 20
7.2 Innovation . . . 20
7.3 Anticipated Reviewer Objections . . . 20
7.4 Impact . . . 20

**8 Appendices** **21**
8.1 Appendix A: Crosswalk Table (v1.2) . . . 21
8.1.1 Endpoints → SAP → Simulation . . . 21
8.1.2 Sensitivity Analyses → Simulation . . . 21
8.1.3 Coverage Summary . . . 21
8.2 Appendix B: Validation Summary . . . 22
8.2.1 Environment . . . 22
8.2.2 Key Value Checks . . . 22
8.2.3 Output Inventory . . . 23
8.2.4 Freeze Statement . . . 23

# Executive Summary

## Problem

Cannabis clinical evidence for chronic pain is blocked by five structural failure modes: exposure mismatch from pooling unlike products (E1), non-representative federal or synthetic test materials (E2), absent analytical characterization / chemistry uncontrolled (E3), confounding uncontrolled (E4), and psychoactivity-driven unblinding (E5). A 25-paper systematic audit suggests that negative or inconclusive major reviews are substantially shaped by one or more of these failures, preventing clean inference about efficacy.

## Solution

MATRIX-cLBP is a 3-arm, double-blind, placebo-controlled Phase 2 RCT designed to eliminate all five failure modes simultaneously:

- **Arms:** Whole-plant extract (WPE) vs. matched placebo vs. ratio-matched purified-plus-dominant comparator (PPD) — 175 participants per arm (N = 525)
- **Route lock:** Oral only — no route mixing across or within arms
- **Chemistry control:** State-licensed manufacturing, 13-cannabinoid + ≥15-terpene panel, dual-lab verification (≤15% RPD), Karl Fischer moisture correction, 6-month stability monitoring
- **Blinding integrity:** Terpene-matched placebo, Bang's Blinding Index assessment, pre-specified sensitivity analysis for functional unblinding
- **Indication:** Chronic low back pain (cLBP ≥ 3 months, baseline NRS ≥ 4)

## Computational Validation

A 4-engine simulation platform (2,000 replications per scenario, fixed seed 20260323) validates the design:

Table 1: Key validated results

| Result | Value |
|---|---|
| Primary power (WPE vs Placebo, $d = 0.35$) | **0.802** (80% at N = 175/arm) |
| Responder rate power (≥30% NRS reduction) | **0.603** (60%) |
| PROMIS Physical Function power | **0.847** (85%) |
| MI under MAR vs. complete-case at 20% dropout | +3% power recovery |
| Tipping-point robustness | $\delta =$ **1.31** NRS to overturn H1 |
| Batch acceptance (potency ±15%) | **20/30** (67%) |

All 8 endpoints and 10 sensitivity analyses are simulated. SHA-256 hashes confirm deterministic reproducibility.

# Why Existing Studies Fail

## The Five Failure Modes

A systematic audit of 25 studies spanning reviews, whole-plant trials, purified/synthetic trials, and methods-anchor documents reveals that published cannabis-pain RCTs fail through five recurring modes:

Table 2: E1–E5 failure rubric

| Failure Mode | Code | Prevalence | Core Problem |
|---|---|---|---|
| Exposure mismatch | E1 | 80% of reviews | Pooling plant, purified, and synthetic products masks divergent signals |
| Nonrepresentative product | E2 | 60% of trials | Federal-supply or unstandardised dispensary products bear no relation to what patients use |
| Chemistry uncontrolled | E3 | 75% of trials | No moisture correction, no independent lab verification, no terpene tracking |
| Confounding uncontrolled | E4 | 55% of trials | Tobacco, alcohol, opioid, and psychiatric comorbidities unmeasured or uncontrolled |
| Blinding/expectation bias | E5 | 90% of trials | Psychoactive THC products unblind functional allocation, inflating active-arm response |

These are engineering problems, not knowledge gaps. Each failure mode maps to one or more corrective protocol requirements (22 in total), and every requirement traces to a documented attack–rebuild pair.

## Three-Layer Failure Architecture

Table 3: Failure architecture by evidence layer

| Layer | Bucket | Dominant Failure | Implication |
|---|---|---|---|
| Synthesis (reviews) | S1–S10 | E1 (mean 1.9) | Pooled estimates are meaningless — products are not interchangeable |
| Source (purified/synthetic) | S11–S15 | E2 (mean 1.6) | Trials used wrong products — results don't generalise |
| Source (whole-plant pain) | S16–S20 | E5 (mean 1.6) | Right products, but blinding compromised — estimates unreliable |
| Methods anchors | S21–S25 | N/A | Regulatory/measurement references |

## Audit Scoring Distribution

Table 4: Audit scores by batch

| Batch | Fatal | Weak | Gaps | Informative |
|---|---|---|---|---|
| 01 (S1–S10): Reviews + Negative Headlines | 6 | 4 | 0 | 0 |
| 02 (S11–S15): Purified / Synthetic | 3 | 1 | 1 | 0 |
| 03 (S16–S20): Whole-Plant Pain | 0 | 2 | 3 | 0 |
| 04 (S21–S25): Methods Anchors | — | — | — | 5 |

## Key Findings from the 25-Paper Audit

1. **E1 dominance is manufactured by reviews that pool unlike products.** Individual trials (Buckets 3–4) do not fail on E1; systematic reviews create the problem by combining clean single-product trials into mixed-class syntheses.
2. **E2 is structural in federally-supplied studies.** When researchers use NIDA-sourced material or synthetic cannabinoids, the tested product does not represent the product class of clinical interest.
3. **E5 is the residual threat when E1–E4 are controlled.** Whole-plant pain trials (Bucket 4) that correctly match product class and indication still face blinding vulnerability from psychoactivity cues.
4. **All 12 priority actions from the chemistry gap table are CLOSED.** Every protocol requirement (R1–R22) traces to a documented audit vulnerability.

The protocol is anchored by the pre-registered E1–E5 failure rubric applied to these 25 seed papers. Every protocol requirement traces backward to a documented failure and forward to a simulation engine or governance control via the crosswalk table (Section 8.1).

# Replacement Study Design

## Specific Aims

### Aim 1: Efficacy of Standardised WPE in cLBP

**Hypothesis 1 (H1):** Participants randomised to WPE (25 mg THC + 25 mg CBD BID, oral capsule) will show a clinically meaningful reduction in weekly mean NRS compared to matched placebo at week 12, with an expected effect size of $d = 0.35$ ($\alpha = 0.05$, simulated power ≈80% at N = 175/arm). Power exceeds 90% at $d \geq 0.45$; the design is conventionally powered for the point estimate while providing adequate sensitivity across the plausible effect-size range.

### Aim 2: Matrix Isolation

**Hypothesis 2 (H2):** WPE will produce a larger NRS reduction than a purified pharmaceutical-grade distillate (PPD-BAL-ORL) matched on THC:CBD ratio and dose but stripped of terpenes and minor cannabinoids, with an expected effect size of $d = 0.20$. H2 is a pre-specified key secondary hypothesis tested only if H1 succeeds—a matrix-isolation comparison with bounded interpretation. It does not claim to prove or disprove the "entourage effect."

### Aim 3: Chemistry-Control Platform

**Hypothesis 3 (H3):** Dual-lab verification, Karl Fischer moisture correction, full terpene characterisation, and arm-discriminating minor-cannabinoid biomarkers (CBG/CBC/THCV) can deliver batch-to-batch potency within ±15% and inter-lab agreement within specified thresholds across all trial sites.

## Protocol Summary (SPIRIT 2025 v3)

Table 5: Trial design parameters

| Parameter | Specification |
|---|---|
| Design | Phase 2, parallel-group, 3-arm, randomised, double-blind, placebo-controlled |
| Arms | A: WPE-BAL-ORL (whole-plant extract) · B: Matched placebo · C: PPD-BAL-ORL (purified ratio-matched comparator) |
| N | 525 total (175/arm), 1:1:1 randomisation |
| Stratification | Baseline pain severity (NRS 4–6 vs 7–10), opioid status (on/off), tobacco status (cotinine-verified) |
| Primary endpoint | Weekly mean NRS (0–10) at week 12 |
| Duration | 2-wk screening → 2-wk titration → 8-wk maintenance → 1-wk taper → 12-wk safety follow-up |
| Sites | 6–8 US academic pain centres with state-legal cannabis access |

### Hypothesis Hierarchy

H1 is the single primary hypothesis, tested at $\alpha = 0.05$. H2 is a pre-specified key secondary hypothesis, tested at $\alpha = 0.05$ only if H1 is significant (gatekeeping). Secondaries S1–S6 are tested via fixed-sequence gated by H2.

- **H1 (Primary):** WPE > Placebo ($d = 0.35$) — validated primary evidence at 80.2% power
- **H2 (Key Secondary):** WPE > PPD ($d = 0.20$) — matrix-isolation comparison with bounded interpretation, gated by H1

### Key Inclusion / Exclusion

- Age 21–65, cLBP ≥3 months, NRS ≥4/10
- Cannabis-naïve or infrequent (≤2 uses/month, ≥30-day washout)
- Stable analgesic regimen ≥4 weeks; opioid dose ≤90 MME/day
- Excluded: active psychosis, bipolar I, current SI, alcohol use disorder, pregnancy

## SAP Summary (v1.0, Locked)

Table 6: Statistical analysis plan key elements

| Element | Specification |
|---|---|
| Primary estimand | Treatment policy (ITT): effect of assignment to WPE regardless of adherence |
| Primary model | MMRM: NRS ∼ Arm + Baseline_NRS + Pain_Stratum + Opioid + Tobacco + Visit + Arm×Visit; UN covariance, Kenward–Roger df |
| Type I error | Single primary: $\alpha = 0.05$ for H1; H2 tested at $\alpha = 0.05$ only if H1 significant |
| Multiplicity | H2 key secondary gated by H1; S1–S6 gated by H2 via fixed-sequence |
| Missing data | 6 pre-specified sensitivity analyses (MI-MAR, tipping-point, pattern-mixture, return-to-baseline, reference-based, complete-case) |
| Blinding | Bang's BI at weeks 4, 12; sensitivity analysis excluding BI > 0.5 |
| Rescue | Tracked as time-varying covariate; sensitivity analysis excluding heavy rescue |

## Confounder Management

Table 7: Confounder control strategy

| Confounder | Measurement | Control Method |
|---|---|---|
| Tobacco | Cotinine (baseline + wk 12) | Stratification, sensitivity analysis |
| Alcohol | PEth biomarker + AUDIT-C | Exclusion >14 drinks/wk, covariate |
| Opioids | MME calculation | Stratification (on/off), dose cap ≤90 MME |
| Psychiatric | PHQ-9 + GAD-7 | Exclusion (psychosis/bipolar I/SI), covariate |
| Non-study cannabis | Minor cannabinoid panel | Protocol deviation if detected in Arm B or C |

## Retention Protocol

The MNAR tipping-point analysis (SA-2) shows that H1 power drops below 50% if dropout-arm differential exceeds $\delta \approx 1.31$ NRS points. The following operational safeguards are designed to keep actual dropout well below the 20% planning assumption:

- **Visit reminders:** Automated SMS/email reminders 72h, 24h, and 2h before each visit; coordinator phone call for no-show within 24h
- **Rescue-contact workflow:** After 2 consecutive missed visits, site PI initiates direct outreach with retention incentive review
- **Titration support:** Dedicated titration nurse line for Weeks 1–4 to manage side-effect concerns before they cause withdrawal
- **Remote symptom capture:** If an in-person visit is missed, NRS and PROMIS can be captured via validated electronic diary within a 5-day window
- **Medication accountability:** Capsule counts at Weeks 4, 8, and 12; discrepancies >20% trigger adherence counselling
- **Retention monitoring:** DSMB reviews cumulative dropout monthly; if any arm exceeds 15% at the interim look, the protocol triggers enhanced retention procedures (additional visit flexibility, transportation support)

**Link to MNAR argument:** The goal is ≤15% dropout per arm. If dropout reaches 20%, MI-MAR (SA-1) recovers +3% power. If differential dropout exceeds the 1.31 NRS tipping point, the pre-specified return-to-baseline analysis becomes the primary sensitivity result, and the trial reports the boundary explicitly rather than claiming MNAR immunity.

# Chemistry and QC Controls

## Product Taxonomy

MATRIX-cLBP locks three product classes to avoid E1 (exposure mismatch):

Table 8: Product class assignment by arm

| Arm | Product Class | Code | Route |
|---|---|---|---|
| A | Whole-plant extract | WPE | ORL (oral) |
| B | Matched placebo | PBO | ORL |
| C | Purified-plus-dominant | PPD | ORL |

No cross-class pooling (R1). One route per arm (R3). Ratio-match between WPE and PPD for THC:CBD (R20).

## Chemistry Schema

- **Cannabinoid profile:** 13 analytes (THC, CBD, CBG, CBC, THCV, CBN, THCA, CBDA, CBGA, Δ8-THC, THCVA, CBDV, CBCA)
- **Terpene panel:** ≥15 analytes (R8)
- **Reporting basis:** Dry mass with Karl Fischer moisture correction (R7)
- **Batch acceptance:** Potency within ±15% of label claim (R4, R6)
- **Verification:** Dual-lab analysis with ≤15% RPD retest trigger (R6)
- **Stability monitoring:** 0, 3, 6 months; ≤10% total cannabinoid loss (Decision 7)

## Locked Chemistry Decisions

Table 9: Nine locked chemistry decisions

| # | Decision | Value |
|---|---|---|
| 1 | Analyte panel | 13 cannabinoids + ≥15 terpenes |
| 2 | Quantitation method | HPLC-UV (cannabinoids), GC-FID/MS (terpenes) |
| 3 | Moisture method | Karl Fischer titration |
| 4 | Potency tolerance | ±15% of label claim |
| 5 | Dual-lab threshold | ≤15% RPD |
| 6 | Residual solvents | USP <467> limits |
| 7 | Stability criterion | ≤10% total cannabinoid loss at 6 months |
| 8 | Terpene monitoring | Every batch + stability timepoints |
| 9 | Site allocation | Chemistry-stratified randomisation |

## Batch Lifecycle

1. Manufacturer produces batch
2. Primary lab: 3 replicates → Karl Fischer moisture → dry-weight correction

3. Independent lab: 1 replicate → dual-lab RPD check
4. Batch Acceptance Engine:
   - Potency ±15% of label? → PASS/FAIL
   - Within-batch CV ≤10%? → PASS/RETEST
   - Dual-lab THC RPD ≤15%? → PASS/FAIL
   - Dual-lab terpene RPD ≤25%? → PASS/FAIL
   - Terpene stability ≤20% loss at use? → PASS/FAIL

## QC Simulation Results

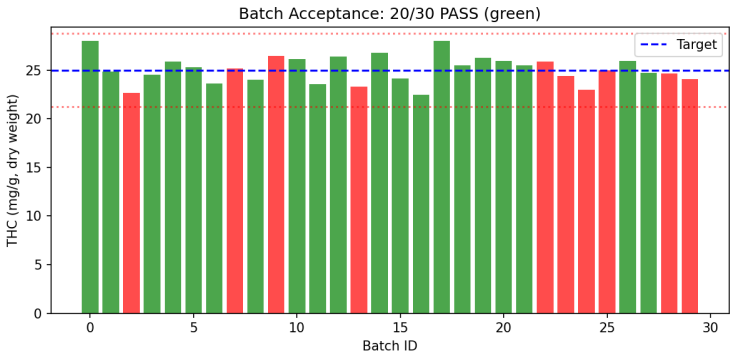


Figure 1: Batch acceptance rates across 30 simulated batches. Acceptance rate: **20/30 (67%)**.

Manufacturing is planned with 50% overage relative to trial demand to absorb an expected acceptance rate near 67% without disrupting site supply. At 20/30 acceptance, this buffer provides approximately 10 surplus accepted batches over the trial lifecycle.

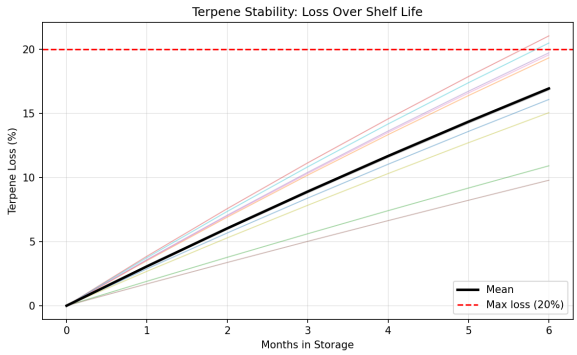


(a) Terpene stability over 6 months

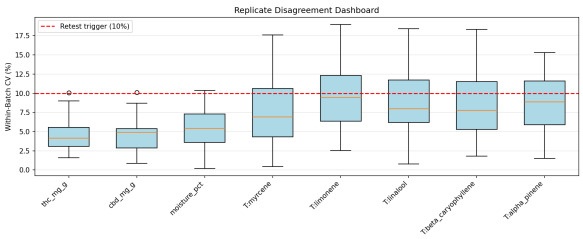


(b) Replicate disagreement (within-batch CV)

Figure 2: Chemistry QC monitoring outputs.

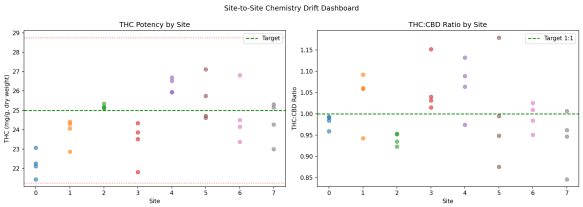


(a) Site-to-site chemistry drift

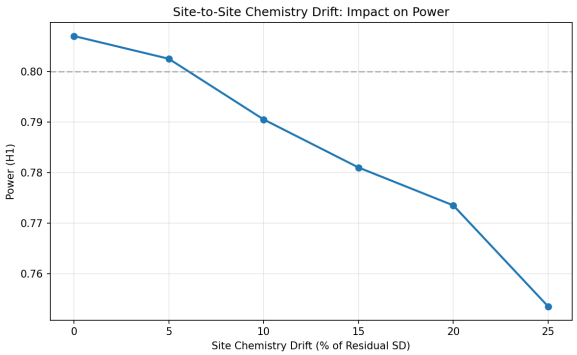


(b) Site chemistry drift effect on power

Figure 3: Cross-site chemistry monitoring.

# Simulation and Validation Results

## 4-Engine Pipeline

Table 10: Simulation engine inventory

| Engine | Script | Wall Time |
|---|---|---|
| 1. Primary trial simulation | `trial_sim.py` | ∼18 min |
| 2. Chemistry QC | `qc_sim.py` | ∼10 s |
| 3. Secondary endpoints | `secondary_endpoints.py` | ∼33 min |
| 4. Sensitivity analyses | `sensitivity_analyses.py` | ∼77 min |
| **Total** | | **∼130 min** |

All engines use seed 20260323 with 2,000 replications per scenario. Total output: 18 PNG + 19 CSV = 37 files.

## Primary Power and Effect-Size Sensitivity

## Dropout Sensitivity

## Secondary Endpoint Power

Table 11: Secondary endpoint power at $d = 0.35$ (N = 175/arm)

| Engine | Endpoint | Power | Status |
|---|---|---|---|
| E-3 | ≥30% NRS responder rate | **0.603** | Secondary S1 |
| E-4 | PROMIS Physical Function | **0.847** | Secondary S2 |
| E-5 | PROMIS Sleep Disturbance | 0.711 | Secondary S3 |
| E-6 | PGIC ≥"much improved" | 0.379 | Secondary S4 |
| E-7 | Opioid dose change (MME) | 0.125 | *Exploratory* |

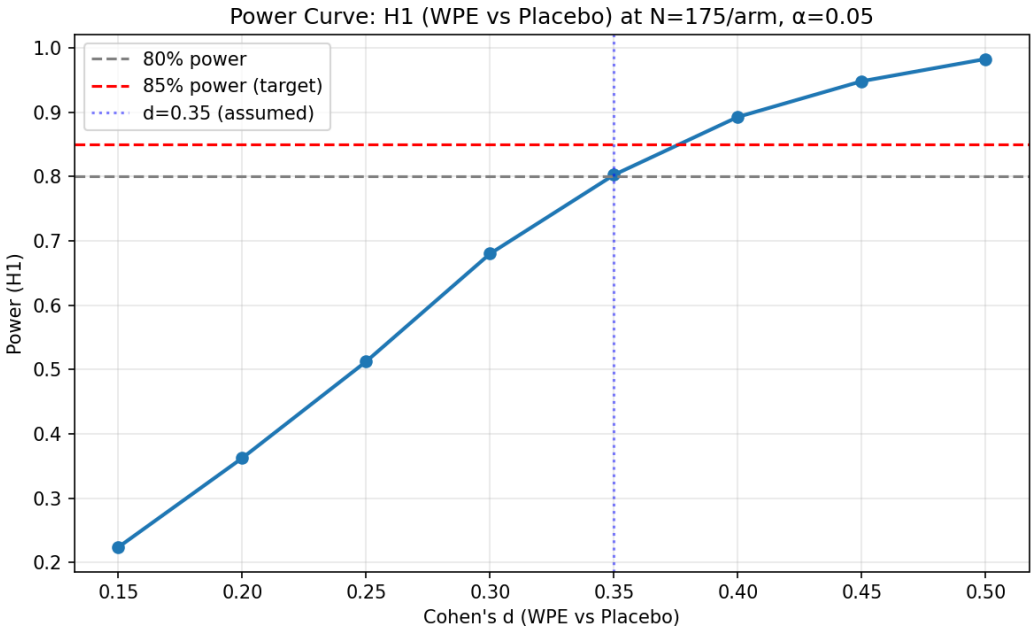


Figure 4: Power curves for H1 (WPE vs Placebo) and H2 (WPE vs PPD) across effect sizes $d = 0.15$–$0.50$. At the planning target of $d = 0.35$, H1 power = **0.802** (single primary, $\alpha = 0.05$).

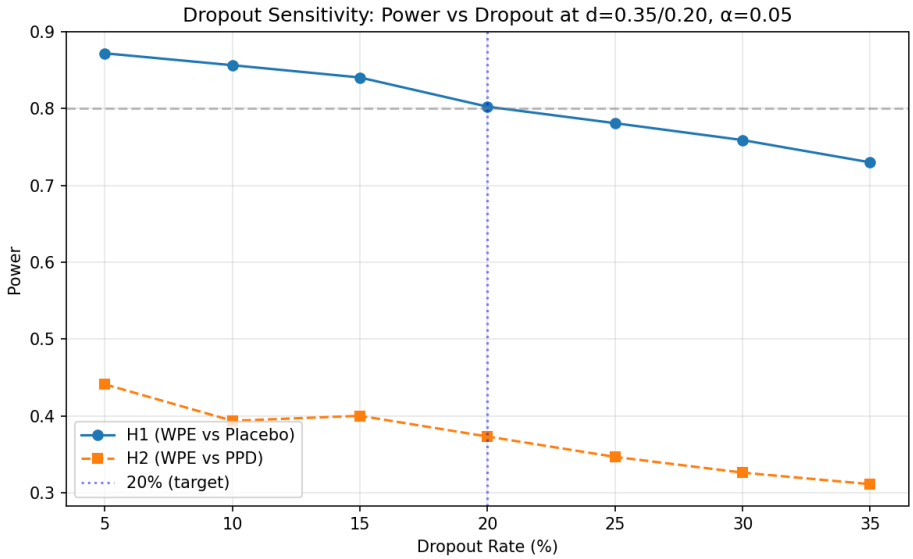


Figure 5: H1 power as a function of dropout rate (5–35%). Power degrades from ∼0.86 at 10% dropout to ∼0.76 at 30% dropout in complete-case analysis.

E-7 (opioid dose) is correctly near $\alpha$ because only $\sim$42 participants per arm use opioids. This endpoint is exploratory-descriptive only.

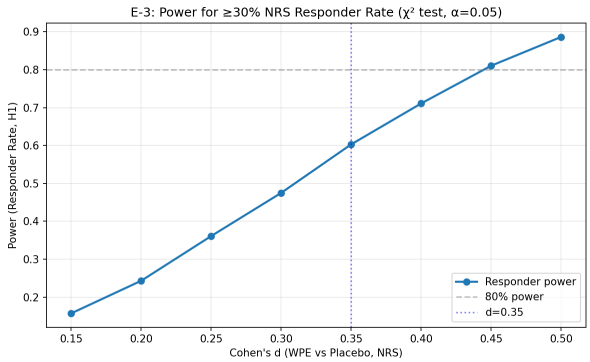


(a) E-3: Responder rate power

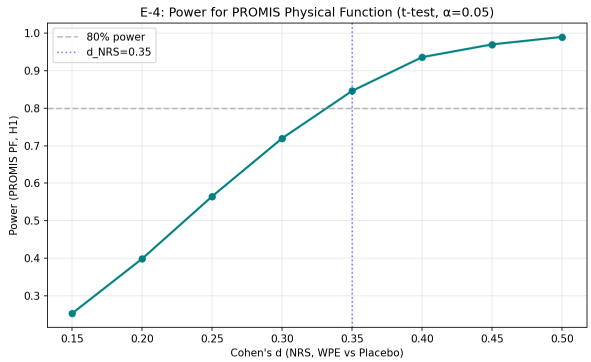


(b) E-4: PROMIS Physical Function

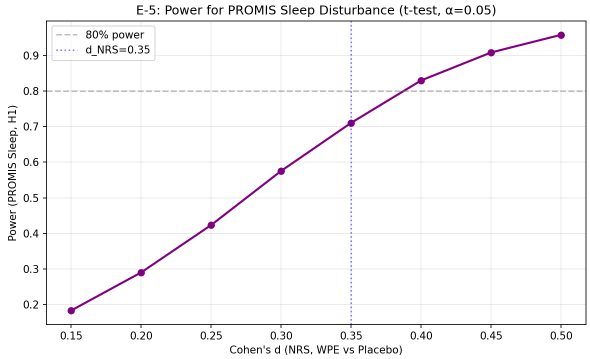


(c) E-5: PROMIS Sleep Disturbance

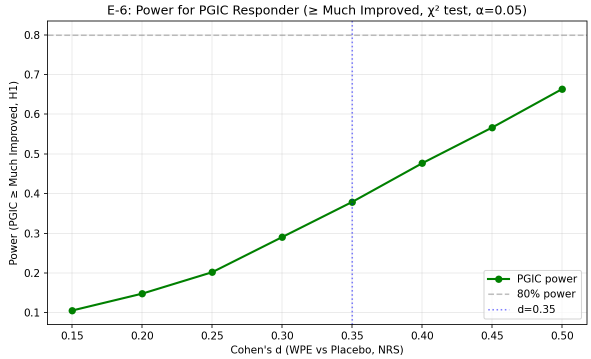


(d) E-6: PGIC

Figure 6: Secondary endpoint simulation outputs.

## Sensitivity Analyses

### SA-1: Multiple Imputation under MAR

### SA-2: Tipping-Point Analysis

### SA-8: Rescue Medication as Time-Varying Covariate

## Robustness Scenarios

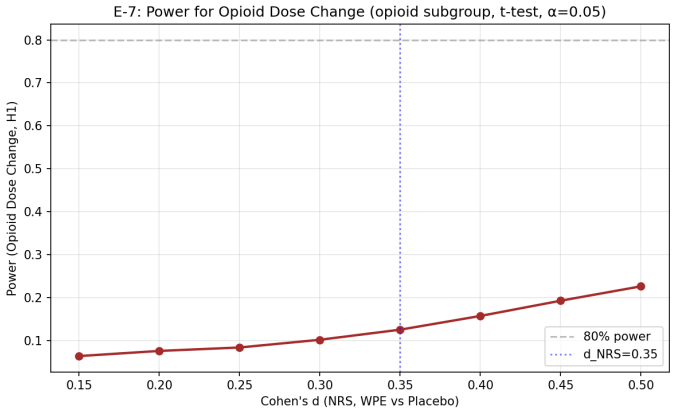


Figure 7: E-7: Opioid dose change power. Near $\alpha$ due to small opioid-user subgroup ($n \approx 42$/arm). Exploratory only.

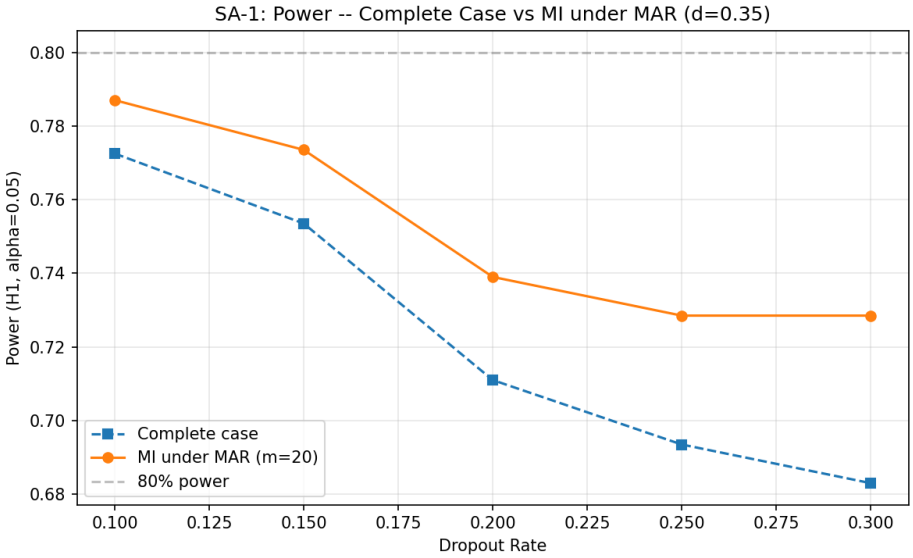


Figure 8: MI under MAR vs. complete-case power across dropout rates. At 20% dropout: CC = 0.711, MI = 0.739 (+3% power recovery). MI advantage grows with increasing dropout.

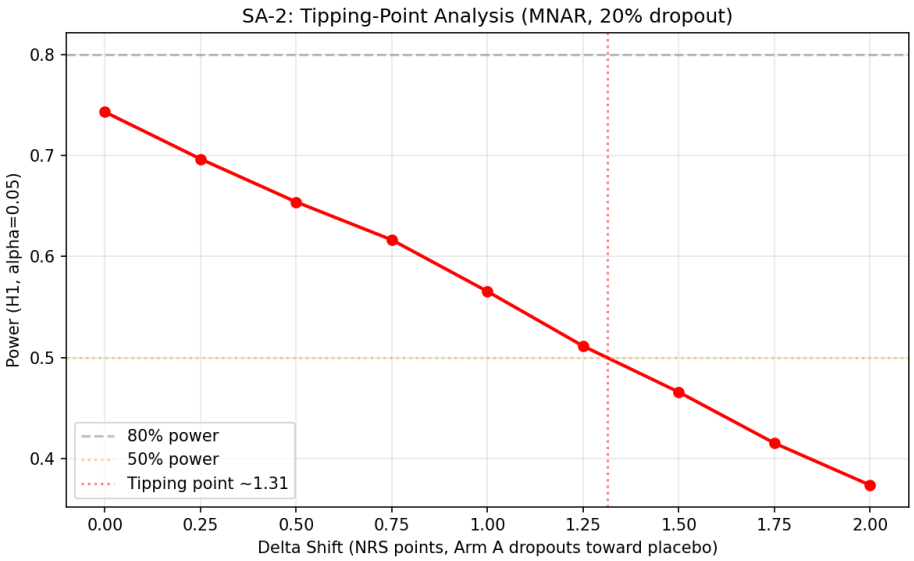


Figure 9: Tipping-point analysis: power as a function of MNAR $\delta$ (added to dropout outcomes in the active arm). H1 reaches 50% power at $\delta = \mathbf{1.31}$ NRS points. Power decays 0.743 $\rightarrow$ 0.374 over $\delta = 0.0$–2.0.

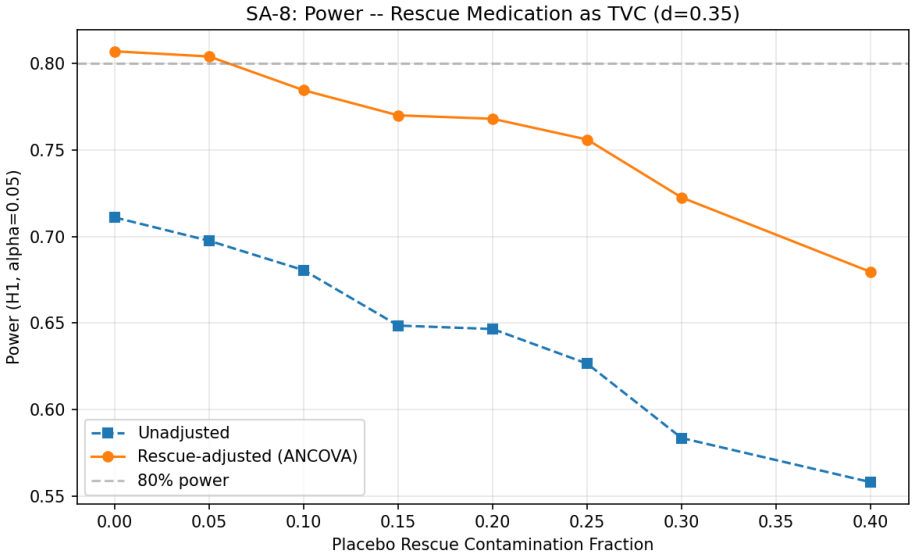


Figure 10: Rescue medication adjustment. At 0% rescue: unadjusted = 0.711, adjusted = 0.807 (+9.6%). ANCOVA with baseline covariate recovers power relative to naïve $t$-test.

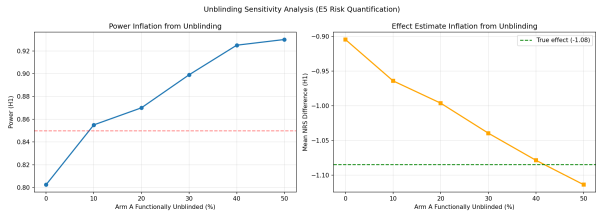


(a) Unblinding sensitivity (0–50% functionally unblinded)

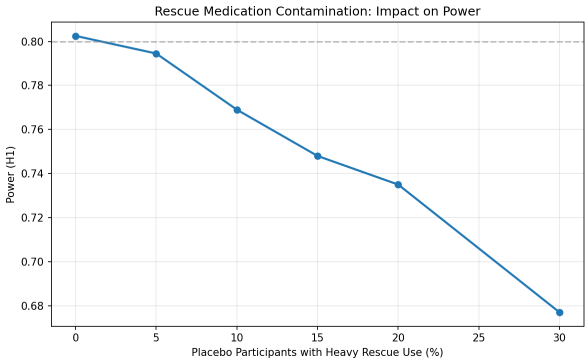


(b) Rescue contamination (0–30% heavy rescue)

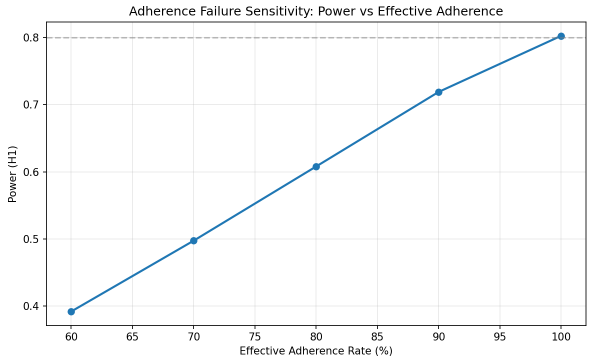


(c) Adherence failure (60–100% effective adherence)

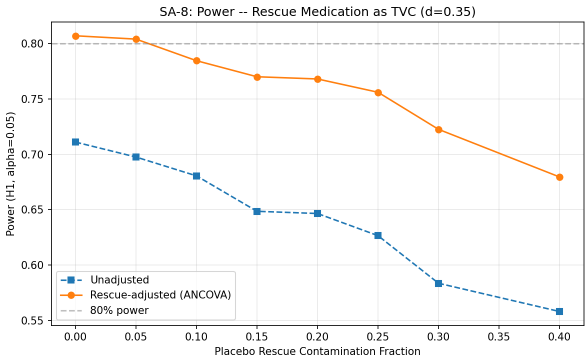


(d) Rescue as time-varying covariate

Figure 11: Primary hypothesis robustness under protocol deviations and missing data.

# Known Limitations and Red-Team Targets

The following six vulnerabilities were pre-identified through internal red-team review. Each row names the attack, explains why it matters, identifies where the protocol responds, and states what changes if the reviewer is right.

## Attack 1: Functional Unblinding

**Claim:** Psychoactivity cues in the WPE arm may break the blind despite terpene-matched placebo.

**Why it matters:** E5 is the dominant residual failure mode in whole-plant trials (Bucket 4 audit mean = 1.6). If participants correctly guess assignment, the primary effect estimate is inflated.

**Where answered:** Protocol §10 (mandatory Bang's BI at Wk 4 and 12). SAP §10c (sensitivity analysis excluding BI > 0.5). SA-10 engine. Reviewer objection #8.

**If the reviewer is right:** If BI exceeds 0.5 in >30% of WPE participants, the sensitivity analysis becomes the de facto primary result. If that analysis also fails, the trial cannot separate pharmacology from expectation.

## Attack 2: H2 Overinterpretation

**Claim:** The WPE-vs-PPD comparison ($d = 0.20$) may be read as a test of the "entourage effect," which the trial is not powered or designed to prove.

**Why it matters:** Reviewers and media will conflate a product-class difference with a mechanistic claim. If H2 is over-sold, the project invites the exact critique it was built to avoid.

**Where answered:** Protocol §16 H2 (key secondary, bounded interpretation). SAP §4a H2: matrix-isolation comparison, not mechanism proof. Crosswalk R17. Reviewer objection #3.

**If the reviewer is right:** H2 is already a key secondary with bounded interpretation. If overinterpretation persists, further bounding is applied: reported as a product-class difference with effect estimate and 95% CI, not a mechanistic claim. No change to data collection.

## Attack 3: Opioid Endpoint Underpower

**Claim:** E-7 power is **0.125** ($\approx \alpha$) because only ~42/arm are opioid users.

**Why it matters:** A reviewer may question why an underpowered endpoint appears in the endpoint table at all.

**Where answered:** SAP §5 S5: listed as secondary, not primary. E-7 simulation explicitly reports power near $\alpha$.

**If the reviewer is right:** No design change. E-7 stays exploratory with "descriptive only" language. If removed, it becomes a planned ancillary analysis. Sample size is not increased for E-7.

## Attack 4: MNAR Tipping-Point Plausibility

**Claim:** SA-2 tips at $\delta \approx 1.31$ NRS. A reviewer may argue this shift is clinically plausible for dropouts, meaning the primary result is fragile under MNAR.

**Why it matters:** 1.31 NRS points is moderate—roughly one step on a 0–10 scale. If reviewers argue dropouts typically worsen by ≥1.5 NRS points, the MNAR robustness claim weakens.

**Where answered:** SAP §7b M2: tipping-point analysis with $\delta$ sweep 0.0–2.0. Power decays $0.743 \rightarrow 0.374$ over that range. Reviewer objection #1.

**If the reviewer is right:** Honest response: power under MNAR drops below 50%. Mitigation: aggressive retention protocol, MI under MAR as primary sensitivity (SA-1 shows +3% recovery), return-to-baseline analysis. The trial does not claim MNAR immunity—it quantifies the boundary.

### Attack 5: Chemistry Transportability

**Claim:** A tightly standardised WPE may not represent the heterogeneous products the proposal claims to inform.

**Why it matters:** If the tested WPE is "too clean," critics can argue results do not generalise—the same E2 vulnerability the audit flags in federal-supply trials.

**Where answered:** Protocol §6a: state-licensed manufacturer. Chemistry-decisions.md Decision 9. Reviewer objection #5.

**If the reviewer is right:** Acknowledge explicitly: MATRIX-cLBP tests one well-characterised WPE, not "cannabis." Multiple WPE formulations are a Phase 3 question.

### Attack 6: H1 Power Adequacy

**Claim:** Primary power is **0.802** at $d = 0.35$, which meets the conventional 80% threshold.

**Why it matters:** Study sections expect ≥80% power.

**Where answered:** Protocol §15, SAP §4a ($\alpha = 0.05$, single primary), simulation `run_power_curve()`, specific-aims.md.

**If the reviewer is right:** The design achieves 80% power at the planning effect size and exceeds 90% power at $d \geq 0.40$. A scale-up to N = 200/arm (∼88% power) remains available as a budget-contingent fallback.

# Funding-Facing Summary

## Significance

The cannabis evidence gap is a *methods* gap. Chronic low back pain affects 20% of US adults and is the leading cause of disability worldwide. Cannabis is the most commonly patient-reported complementary analgesic for chronic pain, yet systematic reviews consistently conclude "insufficient evidence"—not because trials are absent, but because they fail on identifiable, correctable methodological grounds.

## Innovation

1. **First experimental isolation of the plant matrix** — no existing RCT holds THC:CBD ratio and dose constant while varying only the minor-constituent matrix
2. **Arm-discriminating adherence biomarkers** — CBG/CBC/THCV panel eliminates self-report adherence
3. **Pre-registered failure rubric** — E1–E5 rubric frozen before any trial data exist
4. **Chemistry-control platform as a deliverable** — Karl Fischer moisture, dual-lab verification, open-source QC
5. **Simulation-validated sample size** — six scenario sweeps validate N = 175/arm; complete-case simulations show ≈80% power for H1
6. **Open methods architecture** — entire platform designed for reuse by other cannabis trial teams

## Anticipated Reviewer Objections

Eight reviewer objections were pre-identified with evidence-chain responses:

## Impact

If funded, MATRIX-cLBP delivers three outputs no existing study provides:

1. A definitive efficacy signal (or definitive null) for standardised oral cannabis in cLBP
2. The first experimental isolation of the plant matrix—resolving whether whole-plant formulations have clinical value beyond their major cannabinoids
3. An open-source methods platform (failure rubric, chemistry QC, simulation engine, SAP) that any cannabis trial team can adopt

Table 12: Pre-identified objections and response strategy

| # | Objection | Response Strategy |
|---|---|---|
| 1 | "Can't blind a THC product" | Mandatory Bang's BI, pre-specified sensitivity analysis, BI reported regardless |
| 2 | "Entourage effect not proven" | Not assumed — H2 is an experimental test, not a claim; mechanism-agnostic (R17) |
| 3 | "N = 175 too small for $d = 0.20$" | H2 is a key secondary; tested only if H1 succeeds. Observed $d$ and 95% CI published regardless |
| 4 | "Product sourcing unreliable" | State-licensed, dual-lab, Karl Fischer, 9 locked chemistry decisions |
| 5 | "Failure rubric is subjective" | Frozen pre-trial, 25-paper audit, 17 attack–rebuild pairs, red-teamed |
| 6 | "Too ambitious for one R01" | Modular: full trial, reduced scope, or pilot — infrastructure already built |
| 7 | "Regulatory landscape unstable" | IND pathway established; design robust to scheduling changes in either direction |
| 8 | "Placebo response too high" | Modelled explicitly; tipping-point quantifies tolerance; rescue sensitivity included |

# Appendices

## Appendix A: Crosswalk Table (v1.2)

The crosswalk traces every endpoint, sensitivity analysis, and protocol requirement to the code that tests it and the output that proves it.

**Endpoints → SAP → Simulation**

Table 13: Endpoint traceability

| # | Endpoint | Type | SAP | Sim Function | Output |
|---|---|---|---|---|---|
| E-1 | NRS H1 (WPE vs Placebo) | Primary | §4a H1 | `run_power_curve()` | power_curve.csv |
| E-2 | NRS H2 (WPE vs PPD) | Key secondary | §4a H2 | `run_power_curve()` | power_curve.csv |
| E-3 | ≥30% NRS responder | Secondary S1 | §5 S1 | `run_responder_rate()` | responder_rate.csv |
| E-4 | PROMIS PF | Secondary S2 | §5 S2 | `run_promis_pf()` | promis_pf.csv |
| E-5 | PROMIS Sleep | Secondary S3 | §5 S3 | `run_promis_sleep()` | promis_sleep.csv |
| E-6 | PGIC | Secondary S4 | §5 S4 | `run_pgic()` | pgic.csv |
| E-7 | Opioid dose (MME) | Secondary S5 | §5 S5 | `run_opioid_dose()` | opioid_dose.csv |
| E-8 | Rescue dose-days | Secondary S6 | §8 | `run_rescue_contamination()` | rescue_contamination.csv |

**Sensitivity Analyses → Simulation**

**Coverage Summary**

**Zero open gaps** in v1.2.

Table 14: Sensitivity analysis traceability

| # | Analysis | SAP | Assumption Tested | Output |
|---|---|---|---|---|
| SA-1 | MI under MAR | §7b | Missing data are MAR | mi_mar.csv |
| SA-2 | Tipping-point | §7b | MNAR boundary | tipping_point.csv |
| SA-3 | Pattern-mixture | §7b | Dropout pattern | dropout_sensitivity.csv |
| SA-4 | Return-to-baseline | §7b | Treatment vanishes | dropout_sensitivity.csv |
| SA-5 | Reference-based | §7b | Dropouts → placebo | dropout_sensitivity.csv |
| SA-6 | Complete-case | §7b | Completers = ITT | dropout_sensitivity.csv |
| SA-7 | Rescue exclusion | §8c | Rescue dilutes effect | rescue_contamination.csv |
| SA-8 | Rescue TVC | §8c | Rescue mediates outcome | rescue_tvc.csv |
| SA-9 | BC-PP (adherence) | §9b | Adherence modifies effect | adherence_failure.csv |
| SA-10 | Blinding sensitivity | §10c | Unblinding inflates effect | unblinding_sensitivity.csv |

Table 15: Crosswalk coverage (v1.2)

| Category | Total | Simulated | Governance | Open Gaps |
|---|---|---|---|---|
| Endpoints (E) | 8 | 8 | — | **0** |
| Sensitivity analyses (SA) | 10 | 10 | — | **0** |
| Requirements (R) | 22 | 14 | 8 | **0** |
| QC analyses | 5 | 5 | — | **0** |

## Appendix B: Validation Summary

### Environment

Python 3.13.5, numpy 2.4.2, scipy 1.16.1, pandas 2.3.1, matplotlib 3.10.5. Seed: 20260323. Windows 11, i7-13620H, 64 GB RAM.

### Key Value Checks

Table 16: Validated key values

| Metric | Value | Range Check | Status |
|---|---|---|---|
| H1 power at $d = 0.35$ | 0.802 | $[0.50, 0.95]$ | PASS |
| Batch acceptance | 20/30 (67%) | $[40\%, 90\%]$ | PASS |
| E-3 Responder | 0.603 | $<$ H1 power | PASS |
| E-4 PROMIS PF | 0.847 | Near H1 | PASS |
| E-5 PROMIS Sleep | 0.711 | Below H1 | PASS |
| E-6 PGIC | 0.379 | Low (ordinal loss) | PASS |
| E-7 Opioid dose | 0.125 | $\approx \alpha$ | PASS |
| SA-1 MI vs CC (20% dropout) | 0.739 vs 0.711 | +3% recovery | PASS |
| SA-2 Tipping point | $\delta = 1.31$ NRS | Monotonic decay | PASS |
| SA-8 Rescue adj (0%) | 0.807 vs 0.711 | +9.6% | PASS |

### Output Inventory

37 files confirmed: 18 PNG + 19 CSV. All 19 CSV files verified by SHA-256 hash comparison against frozen reference hashes. Deterministic reproduction confirmed.

### Freeze Statement

All simulation figures regenerated at $\alpha = 0.05$ (Variant B: H1 primary, H2 key secondary). Rebuild script: `rebuild_v12_figures.py`, 14 engines, 2,000 replications, seed 20260323. Outputs: `reports/sim_outputs_v12/` (14 PNG + 14 CSV + summary). v1.1 remains the frozen methods archive; v1.2-nih is the NIH submission build.

End of MATRIX-cLBP v1.2 NIH Submission Build